

PARTIAL compact-inference latency vs raw dense · NVIDIA RTX A6000 · batch 1
Only 3/3 seeds with sealed full workers and identical paired panels; gray cells are missing
GPU raw → logits

