


PARTIAL compact-inference latency vs raw dense · NVIDIA RTX A6000 · batch 8
Only 3/3 seeds with sealed full workers and identical paired panels; gray cells are missing
GPU raw → logits



Host raw → logits

